

Labs: Trees, Hierarchical Clustering,
Outlier detection, Cook's Distance

Project Dataset One-on-One with the
instructor

Thilanka Munasinghe

Data Analytics

ITWS-4600/ITWS-6600/MATP-4450/CSCI-4960

Group 3 - Lab 6, 21st October 2022

Work on the remaining scripts (Group 3)

- <https://rpi.box.com/s/2xx9ul1fmc6bf5ff8h4jreae69emikmf>
- **What is expected:** You are asked to Explore, Inspect the code/scripts that are on Group 3 in the Rstudio environment and get familiar with those scripts.

Scripts – work through these

Reminder to finish these code

examples See in folder group2/ Lab1

Go over the following scrips,

Lab1_bronx1.R.

Lab1_bronx2.R

Lab1_ctree2.R

Lab1_kknn1.R

Lab1_kknn2.R

Lab1_kknn3.R

Lab1_kmeans1.R

Lab1_nyt.R

Search before you ask! You might need to search your code errors online when you are debugging your code!.

script fragments in R available on:

<https://rpi.box.com/s/2xx9ul1fmc6bf5ff8h4jreae69emikmf>

NOTE: you are allowed to work in small groups and discuss during this lab.

Scripts – work through these

Next...

See in folder group2/ Lab3

Go over the following scrips,

Lab3_ctree1.R

Lab3_ctree2.R

Lab3_ctree3.R

.....

And the remaining code snippets in

group2/Lab 2 and Lab3

Search before you ask! You might need to search your code errors online when you are debugging your code!

script fragments in R available on:

<https://rpi.box.com/s/2xx9ul1fmc6bf5ff8h4jreae69emikmf>

NOTE: you are allowed to work in small groups and discuss during this lab.

Scripts – work through these

Next...

See in folder group2 and group3/

Labs

Go over the following scrips,

Lab3_ctree1.R

Lab3_ctree2.R

Lab3_ctree3.R

.....

And the remaining code snippets in
group2/Lab 2 and Lab3

**Search before you ask! You might need to search your
code errors online when you are debugging your code!**

script fragments in R available on:

<https://rpi.box.com/s/2xx9ul1fmc6bf5ff8h4jreae69emikmf>

**NOTE: you are allowed to work in small groups and
discuss during this lab.**

Outlier Detection

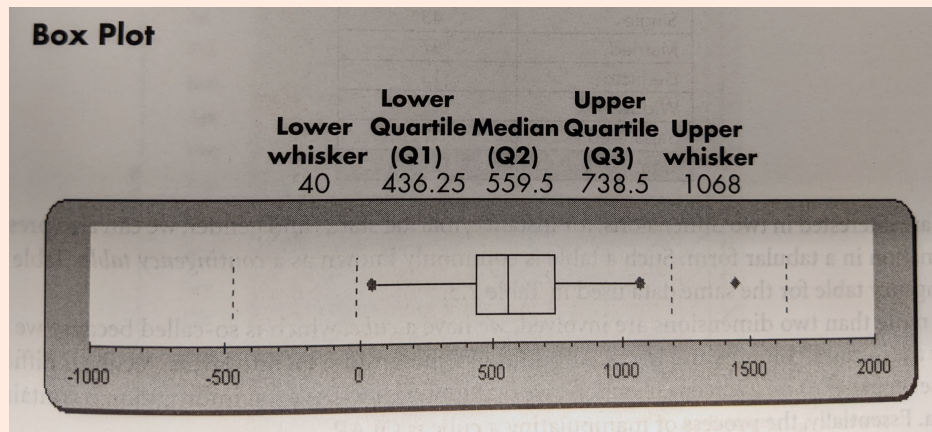
- Outliers are atypical (extreme) values.
- Although they are not errors (sometimes!)
- We need to pay careful attention to outliers.
- Sometimes, it is necessary to identify them and remove from the data and this is when we are looking for typical cases.
- Our objective is to see the performance at the extremes.
- Outliers provide very valuable and useful information.

Using Interquartile Range (IQR) to detect Outliers

- IQR is the measure of the variability, that is more robust than the standard deviation for skewed data.
- IQR is the difference between the $Q3 - Q1$ ($Q1=25^{\text{th}}$ percentile and $Q3=75^{\text{th}}$)
- It is important to check the Interquartile Range (IQR) for your data.
- Check how many data points are outside the IQR. Easiest way to check IQR is by generating boxplots, but this visual inspection is questionable and very subjective from person to person
- So we will use a metric generated using IQR.

Boxplot

- Boxplot is a useful graphical tool for mapping the data behavior in the middle as well as at the end of the distributions.
- Boxplots uses the median (Q2) and the lower quartile and upper quartile (alternative names are)
lower quartile= $Q1=25^{\text{th}}$ percentile
upper quartile = $Q3=75^{\text{th}}$ percentile



Outliers: outside the fence

- Boxplot is constructed by drawing a box between the upper and lower quartiles with a solid line drawn across the box to locate the median value.
- Following quantities are also called “***fences***”.
- “***Fences***” are essential for identifying the extreme values in the tails of the distributions.

Lower Inner Fence: $Q1 - 1.5 \times (IQR)$

Lower Inner Fence: $Q3 + 1.5 \times (IQR)$

Rule of Thumb

- Values are outside the Lower Inner Fence and Uppers Inner Fence can be considered as extreme values.
- Note: existing many outlier points indicates that the dataset may not be suitable for reliable modeling.
- There are other more prominent methods to detect outliers than using IQR lower and inner fencing technique.
- Cook's Distance is also used to identify the outlier points in data.

Cook's Distance

- Cook's distance is a commonly used estimate of the influence of a data point when performing a least-squares regression analysis.
- Cook's distance can be used in several ways:
 - [1] to indicate influential data points that are particularly worth checking for validity;
 - [2] or to indicate regions of the design space where it would be good to be able to obtain more data points.

Cook's Distance

- Each element in the Cook's distance D is the normalized change in the fitted response values due to the deletion of an observation. The Cook's distance of observation i is

$$D_i = \frac{\sum_{j=1}^n (\hat{y}_j - \hat{y}_{j(i)})^2}{p \text{ MSE}},$$

Where:

- \hat{y}_j is the j^{th} fitted response value.
- $\hat{y}_{j(i)}$ is the j^{th} fitted response value, where the fit does not include observation i .
- MSE is the mean squared error.
- p is the number of coefficients in the regression model.

Regression Model using lm()


We can generate the regression model using lm() function,
Our response variable is mile per gallon (mpg)
and predictor variables are number of cylinders (cyl) and weight (wt)

```
# Cook's Distance' example using mtcars
mtcars
head(mtcars)
str(mtcars)
model1 <- lm(mpg ~ cyl + wt , data = mtcars)
model1
help("cooks.distance")
plot(model1, pch = 18, col= 'red', which = c(4))
```

```
# we can use the cooks.distance() function to identify the Cook's distance to
# each observation
cooks.distance(model1)
CooksDistance <- cooks.distance(model1)
```

Interpretation

Cook's Distance visualization for the regression model: $\text{lm}(\text{mpg} \sim \text{cyl} + \text{wt})$



cooks.distance() function

We can use the `cooks.distance()` function to identify the Cook's distance to each observation in our dataset

```
> cooks.distance(model1)
```

Mazda RX4	Mazda RX4 Wag	Datsun 710	Hornet 4 Drive	Hornet Sportabout	Valiant	Duster 360
0.0050772590	0.0004442585	0.0567764620	0.0018029260	0.0235271472	0.0050205614	0.0178733213
Merc 240D	Merc 230	Merc 280	Merc 280C	Merc 450SE	Merc 450SL	Merc 450SLC
0.0091033181	0.0065061176	0.0004643600	0.0075293380	0.0116847953	0.0102875723	0.0005228914
Cadillac Fleetwood	Lincoln Continental	Chrysler Imperial	Fiat 128	Honda Civic	Toyota Corolla	Toyota Corona
0.0035498738	0.0001501537	0.3189363624	0.1592990291	0.0276449872	0.2233281268	0.0913548207
Dodge Challenger	AMC Javelin	Camaro Z28	Pontiac Firebird	Fiat X1-9	Porsche 914-2	Lotus Europa
0.0040263378	0.0120218543	0.0165559199	0.0569730451	0.0001790454	0.0033281614	0.0216355209
Ford Pantera L	Ferrari Dino	Maserati Bora	Volvo 142E			
0.0237336584	0.0105550987	0.0072685192	0.0727399065			

```
# we can use the cooks.distance() function to identify the Cook's distance to  
# each observation  
cooks.distance(model1)  
CooksDistance <- cooks.distance(model1)  
  
# Now we will round off the values to 5 decimal points so that it is easy to read  
# we can use the round() function to round off values in R.  
round(CooksDistance, 5)
```



```
# Now we will round off the values to 5 desimal points so that it is easy to read
# we can use the round() function to round off values in R.
round(CooksDistance, 5)
```

> CooksDistance

Mazda RX4	Mazda RX4 Wag	Datsun 710	Hornet 4 Drive	Hornet Sportabout	Valiant	Duster 360
0.0050772590	0.0004442585	0.0567764620	0.0018029260	0.0235271472	0.0050205614	0.0178733213
Merc 240D	Merc 230	Merc 280	Merc 280C	Merc 450SE	Merc 450SL	Merc 450SLC
0.0091033181	0.0065061176	0.0004643600	0.0075293380	0.0116847953	0.0102875723	0.0005228914
Cadillac Fleetwood	Lincoln Continental	Chrysler Imperial	Fiat 128	Honda Civic	Toyota Corolla	Toyota Corona
0.0035498738	0.0001501537	0.3189363624	0.1592990291	0.0276449872	0.2233281268	0.0913548207
Dodge Challenger	AMC Javelin	Camaro Z28	Pontiac Firebird	Fiat X1-9	Porsche 914-2	Lotus Europa
0.0040263378	0.0120218543	0.0165559199	0.0569730451	0.0001790454	0.0033281614	0.0216355209
Ford Pantera L	Ferrari Dino	Maserati Bora	Volvo 142E			
0.0237336584	0.0105550987	0.0072685192	0.0727399065			

Interpretation

we can sort the values in ascending order
sort(round(CooksDistance, 5))

```
> sort(round(CooksDistance, 5))
```

Lincoln Continental	Fiat X1-9	Mazda RX4 Wag	Merc 280	Merc 450SLC	Hornet 4 Drive	Porsche 914-2
0.00015	0.00018	0.00044	0.00046	0.00052	0.00180	0.00333
Cadillac Fleetwood	Dodge Challenger	Valiant	Mazda RX4	Merc 230	Maserati Bora	Merc 280C
0.00355	0.00403	0.00502	0.00508	0.00651	0.00727	0.00753
Merc 240D	Merc 450SL	Ferrari Dino	Merc 450SE	AMC Javelin	Camaro Z28	Duster 360
0.00910	0.01029	0.01056	0.01168	0.01202	0.01656	0.01787
Lotus Europa	Hornet Sportabout	Ford Pantera L	Honda Civic	Datsun 710	Pontiac Firebird	Volvo 142E
0.02164	0.02353	0.02373	0.02764	0.05678	0.05697	0.07274
Toyota Corona	Fiat 128	Toyota Corolla	Chrysler Imperial			
0.09135	0.15930	0.22333	0.31894			

Outlier Detection: Cook's Distance

Lab / In-Class work:

```
# Outlier Detection using "Cooks Distance"
# Multivariate Regression using Cook's Distance
# Cook's Distance is an estimate of the influence of a data point.
# Cook's Distance is a summary of how much a regression model changes when the ith observation is removed from the data.
library(ISLR)
# Let's look at the baseball hitters dataset in ISLR package.
head(Hitters)
dim(Hitters)
is.na(Hitters) # check for the missing values.
# Now we will remove the NA (missing values) using the na.omit() function
HittersData <- na.omit(Hitters)
dim(HittersData) # checking the dimensions after removing the NAs.
glimpse(HittersData)
head(HittersData)
# Now we will implement a multivariate regression model using all the features in the dataset to
# predict the salary of the Baseball player
SalaryPredictModel1 <- lm(Salary ~., data = HittersData)
summary(SalaryPredictModel1)
# Multiple R-squared:  0.5461, Adjusted R-squared:  0.5106
```

```

# Cook's Distance.
cooksD <- cooks.distance(SalaryPredictModel1)
influential <- cooksD[(cooksD > (3 * mean(cooksD, na.rm = TRUE)))]
influential
# We see that 18 players have a Cook's Distance greater than 3x the mean.
# Let's exclude these 18 players and rerun the model to see if we have a better fit in our regression model.
names_of_influential <- names(influential) # checking the names of the influential/outlier players
names_of_influential
outliers <- HittersData[names_of_influential,]
Hitters_Without_Outliers <- HittersData %>% anti_join(outliers)

# Model 2: without the outliers
SalaryPredictModel2 <- lm(Salary ~. , data = Hitters_Without_Outliers)
summary(SalaryPredictModel2)
# Multiple R-squared:  0.6721, Adjusted R-squared:  0.6445
# We have improved from an Adjusted R-Squared of 0.5106 to 0.6445 with the removal of only 18 observations

```

Trees for the Titanic

Reminder:

`data(Titanic)`

`rpart`, `ctree`, `hclust`, for:

Survived ~ .

Read the titanic dataset documentation in Rdocumentation:

<https://www.rdocumentation.org/packages/titanic/versions/0.1.0>

Feedback: Assignment 5 Presentation

- Make sure to meet with the instructor during today's class One-on-One time to get the Assignment 5 (Project Proposal Presentation) Feedback.
- You **MUST** get the Assignment 5 feedback from the instructor before you continue to work on the project.
- It is important to get the feedback and discuss your project with the instructor.

Assignment 4

Reminder: Assignment 4 is available on LMS.